

EECS 370

Introduction to Computer Organization

The Complete Open-Book Exam Cheat Sheet

Everything from binary to virtual memory — taught from scratch



This guide covers: number representation · the LC2K & ARM/LEGV8 ISAs · assembly & C translation · assemblers, linkers & loaders · calling conventions & the stack · digital logic & FSMs · single-/multi-cycle & pipelined datapaths · hazards & forwarding · performance & Amdahl's Law · caches & AMAT · virtual memory & TLBs.

0 Contents

1	Big Picture: What This Course Is About	2
2	Number Representation & Binary Arithmetic	2
2.1	Positional notation and bases	2
2.2	Unsigned vs. signed (two's complement)	2
2.3	Sign extension	3
2.4	Addition, overflow, and carry	3
3	Instruction Set Architecture (ISA) Fundamentals	3
4	The LC2K ISA (Project 1 core)	4
4.1	Instruction formats and the exact bit layout	4
4.2	The 8 instructions (memorize this table)	4
4.3	Building everything from 8 instructions	4
5	LC2K Assembly, the Assembler, Linker & Loader	5
5.1	Assembly source format	5
5.2	Two-pass assembler	5
5.3	Linker: combining multiple files	5
6	The ARM / LEGv8 ISA	6
6.1	Registers	6
6.2	Core instructions	6
7	C → Assembly, Calling Conventions & the Stack	7
7.1	Translating C control flow	7
7.2	The stack and stack frames	7
7.3	Standard function prologue / epilogue	7
8	Digital Logic	8
8.1	Boolean algebra & gates	8
8.2	Combinational building blocks	8
8.3	Sum-of-products from a truth table	8
8.4	Sequential logic: state & clocks	8
8.5	Finite State Machines (FSM)	8
9	Single-Cycle Datapath	9
9.1	The five things every instruction does	9
10	Multi-Cycle Datapath	9
11	Pipelining	9
11.1	The classic 5-stage pipeline	9
11.2	Hazards — the three reasons pipelines stall	10


11.2.1	Data hazards & forwarding	10
11.2.2	Control hazards & branch handling	10
11.2.3	Dependency types (know the vocabulary)	11
11.2.4	Dynamic branch prediction: the 2-bit saturating counter	11
11.3	The LC2K Pipeline Lab (Project 3) — implementation cheat sheet	11
11.3.1	The pipeline-register and machine state structs	11
11.3.2	What each stage does (and which register it writes)	12
11.3.3	Forwarding (done entirely in EX) — the exact algorithm	12
11.3.4	The load-use stall (the one case forwarding can't fix)	12
11.3.5	Branch handling: predict-not-taken, resolve in MEM	13
11.3.6	Per-cycle driver & start/stop	13
12	Performance Analysis	13
13	Caches & the Memory Hierarchy	14
13.1	Address breakdown	14
13.2	Mapping: direct, set-associative, fully associative	14
13.3	The 3 C's — classify every miss	14
13.4	Replacement & write policies	15
13.5	Performance: AMAT	15
14	Virtual Memory	15
14.1	Paging and address translation	15
14.2	Page table	16
14.3	TLB — making translation fast	16
15	Rapid-Fire Exam Checklists	16

1 Big Picture: What This Course Is About

The One-Sentence Summary

EECS 370 explains **how a program written in C becomes electrons moving through transistors**: $C \rightarrow \text{assembly} \rightarrow \text{machine code} \rightarrow \text{executed by a datapath built from logic gates}$, while a memory hierarchy (registers $\rightarrow$ cache $\rightarrow$ RAM $\rightarrow$ disk via virtual memory) feeds it data fast enough to matter.

The course is a vertical slice through the computing stack. Read this as a pipeline:



How to use this sheet in the exam

Each topic has a **procedure** (numbered steps you can follow mechanically) plus the **formula** and a **worked example**. On a timed exam, find the matching procedure, plug in numbers, and check against the "Common Mistake" boxes.

2 Number Representation & Binary Arithmetic

2.1 Positional notation and bases

A number in base b with digits $d_{n-1} \dots d_1 d_0$ has value $\sum d_i b^i$. We use **binary** (base 2), **hex** (base 16, digits 0–9, A–F), and decimal.

Conversions you must do instantly

- **Bin→Hex**: group bits into 4s from the right; each group is one hex digit.
- **Hex→Bin**: expand each hex digit to 4 bits (A = 1010, F = 1111).
- n **bits** represent 2^n distinct values.
- **Powers**: $2^{10} = 1024 \approx 10^3$ (Ki), 2^{20} (Mi), 2^{30} (Gi).

2.2 Unsigned vs. signed (two's complement)

With n bits:

- **Unsigned** range: 0 to $2^n - 1$.
- **Two's complement (signed)** range: -2^{n-1} to $2^{n-1} - 1$. The MSB has weight -2^{n-1} ; all other bits have their normal positive weight.

$$10110101_2 \begin{array}{c} -2^7 \\ \boxed{1} \end{array} \begin{array}{c} 2^6 \\ \boxed{0} \end{array} \begin{array}{c} 2^5 \\ \boxed{1} \end{array} \begin{array}{c} 2^4 \\ \boxed{1} \end{array} \begin{array}{c} 2^3 \\ \boxed{0} \end{array} \begin{array}{c} 2^2 \\ \boxed{1} \end{array} \begin{array}{c} 2^1 \\ \boxed{0} \end{array} \begin{array}{c} 2^0 \\ \boxed{1} \end{array} = -128 + 32 + 16 + 4 + 1 = -75$$

Negate in two's complement: "flip and add 1"

To compute $-x$: invert every bit (one's complement), then add 1. Example with 8 bits: $x = 5 = 00000101$. Flip $\rightarrow 11111010$, $+1 \rightarrow 11111011 = -5$. **Check:** $x + (-x)$ should be 0 (with the carry-out discarded).

2.3 Sign extension

To widen a signed number to more bits, **copy the sign bit** into all new high bits ($0111 \rightarrow 00000111$; $1011 \rightarrow 11111011$). For unsigned, pad with zeros. This is why the LC2K `lw/sw/beq` offset (16 bits) is *sign-extended* to 32 bits before use.

2.4 Addition, overflow, and carry

Add like decimal but carry at 2. Two's-complement add/subtract use the *same* hardware.

- **Unsigned overflow** = carry-out of the MSB.
- **Signed overflow** = carry-into MSB $\neq$ carry-out of MSB. Equivalently: $(+) + (+) = (-)$ or $(-) + (-) = (+)$. Adding opposite signs *never* overflows.

Common Mistake

"The number is negative" $\neq$ "overflow occurred." Overflow means the *true* result doesn't fit in n bits. Always reason about the sign-bit rule, not just the result's sign.

3 Instruction Set Architecture (ISA) Fundamentals

What an ISA is

The ISA is the **contract between hardware and software**: the set of instructions, the registers, the memory model, and the binary encoding. Software targets the ISA; any hardware that implements the ISA correctly will run that software. It is the boundary that lets one program run on many chip generations.

Key ISA design axes:

- **RISC vs. CISC:** RISC (LC2K, ARM, MIPS) = few, simple, fixed-length instructions, load/store architecture. CISC (x86) = many complex variable-length instructions.
- **Load/store architecture:** ALU operations work only on registers; memory is touched *only* by load (`lw`) and store (`sw`). LC2K and ARM are load/store.
- **Word size:** LC2K uses 32-bit words and is **word-addressed** (address i = the i -th word). Most real machines (ARM) are **byte-addressed**.
- **Endianness:** byte order of a multi-byte word. Little-endian stores the least-significant byte at the lowest address; big-endian the reverse.


4 The LC2K ISA (Project 1 core)

LC2K at a glance


8 general registers (`r0–r7`), each **32 bits**. `r0` is an ordinary register but conventionally holds **0** (the assembler does *not* force it). Memory is an array of **32-bit words**, **word-addressed**. Every instruction is **one 32-bit word**. There are exactly **8 instructions** in **4 formats**.

4.1 Instruction formats and the exact bit layout


Bits are numbered 31 (MSB) down to 0 (LSB). The opcode is always bits **24–22**.



J-type (jalr)



O-type (halt, noop)



4.2 The 8 instructions (memorize this table)

Instr	Opcode	Type	Assembly	Action (RTL)
add	000 (0)	R	add rA rB rDest	reg[Dest] = reg[A] + reg[B]
nor	001 (1)	R	nor rA rB rDest	reg[Dest] = ~(reg[A] reg[B])
lw	010 (2)	I	lw rA rB off	reg[B] = mem[reg[A] + off]
sw	011 (3)	I	sw rA rB off	mem[reg[A] + off] = reg[B]
beq	100 (4)	I	beq rA rB off	if(reg[A]==reg[B]) PC=PC+1+off
jalr	101 (5)	J	jalr rA rB	reg[B]=PC+1; PC=reg[A]
halt	110 (6)	O	halt	increment PC, then stop machine
noop	111 (7)	O	noop	do nothing (PC++)

The two “gotchas” everyone trips on

- beq target is PC-relative and PC is *already incremented*:** target = (address of the beq) + 1 + offset. So offset = (target address) – (beq address) – 1.
- jalr edge case:** if `regA == regB`, the spec says `reg[B]` is first written with `PC+1` and then PC takes the *new* value of that register — read the spec carefully; the link write happens before the jump uses `regA`.

4.3 Building everything from 8 instructions

LC2K is deliberately minimal — you synthesize the rest:

- **Load a constant** (no “addi”): put the value in memory with `.fill` and `lw` it, or build it from `r0`.

- **NOT x:** `nor x x dest` ($\overline{x|x} = \overline{x}$).
- **Move / copy:** `add src r0 dest` (assuming `r0=0`).
- **Subtract:** negate then add. To get $-y$: `nor y y t` then add a known 1 (two's complement), or use `.fill` constants.
- **Unconditional jump:** `beq r0 r0 off` (always taken).
- **AND** via De Morgan: $a \& b = \overline{\overline{a} | \overline{b}}$ using three `nor`s.

5 LC2K Assembly, the Assembler, Linker & Loader

5.1 Assembly source format

Each line: `[label] opcode field0 field1 field2 # comment`. Labels start in column 1 (max 6 chars, begin with a letter). Fields are decimal numbers *or* labels (for the offset/address fields).

```

        lw      0 1 five    // reg1 = mem[five] = 5
        lw      0 2 neg1    // reg2 = -1
start   add      1 2 1      // reg1 = reg1 + reg2 (count down)
        beq     0 1 done    // if reg1==0 goto done (reg0 is 0)
        beq     0 0 start   // unconditional jump back
done    halt
five    .fill    5
neg1    .fill    -1
stk     .fill    0

```

The `.fill` directive

`.fill N` places the 32-bit value `N` into the next memory word. `N` may be a number *or* a label — if it's a label, the assembler stores that label's *address*. This is how you make constants, jump tables, and pointers. `.fill` is *not* an instruction; it's data placed inline in the program image.

5.2 Two-pass assembler

1. **Pass 1 — build the symbol table.** Walk every line, tracking the address (line number, since each line is one word). Record each label $\rightarrow$ address. Catch duplicate labels.
2. **Pass 2 — generate machine code.** For each line, look up any label operands in the symbol table, compute offsets (for `beq`: `labelAddr - (PC+1)`), assemble the 32-bit word, and emit it as a decimal integer (one per line).


Why two passes?

Forward references. A `beq` can jump to a label defined later in the file. You cannot fill in its offset until you know that label's address, so pass 1 finds all addresses first, then pass 2 resolves them.

5.3 Linker: combining multiple files

Real programs are split across files compiled separately. The linker stitches them together.

- **Global labels:** a label is global if it's referenced across files. Two label kinds in object files:
 - **Defined (Symbol Table):** “I provide label `x` at offset `k` in my text/data.”
 - **Undefined (Relocation Table):** “I use label `y` but don't define it — patch me.”
- **Relocation:** when files are concatenated, every absolute address (e.g. a `.fill label` or an `lw/sw` that uses a label) must be *relocated* by adding the base offset of its section. PC-relative `beq` offsets do *not* need relocation (they are self-relative within the file). This is exactly Project 2's linker logic.
- **Loader:** copies the final image into memory at its load address and starts execution at the entry point.



6 The ARM / LEGv8 ISA

Why ARM too?

LC2K teaches the bare minimum; ARM (the LEGv8 teaching subset from Patterson & Hennessy) shows a real, byte-addressed, 64-bit RISC ISA with a richer instruction set and a real calling convention.

6.1 Registers

32 registers, 64 bits each: X0–X30 plus XZR (the zero register, reads 0).

Register	Role	Saved by
X0–X7	arguments / return values	caller-saved
X8	indirect result location	—
X9–X15	temporaries	caller-saved
X16–X17	scratch (IP0/IP1)	caller-saved
X18	platform register	—
X19–X27	saved registers	callee-saved
X28 (SP)	stack pointer	callee-saved
X29 (FP)	frame pointer	callee-saved
X30 (LR)	link register (return address)	callee-saved
XZR	constant zero	—

6.2 Core instructions

Instruction	Meaning	Format
ADD Xd, Xn, Xm	$X_d = X_n + X_m$	R
SUB Xd, Xn, Xm	$X_d = X_n - X_m$	R
ADDI Xd, Xn, #imm	$X_d = X_n + \text{imm}$	I
SUBI Xd, Xn, #imm	$X_d = X_n - \text{imm}$	I
AND/ORR/EOR ...	bitwise and / or / xor	R
LSL/LSR Xd,Xn,#sh	logical shift left / right	R
LDUR Xt, [Xn,#off]	$X_t = \text{Mem}[X_n + \text{off}]$ (load 64-bit)	D
STUR Xt, [Xn,#off]	$\text{Mem}[X_n + \text{off}] = X_t$ (store)	D
CBZ Xt, label	if $X_t == 0$ branch	CB
CBNZ Xt, label	if $X_t \neq 0$ branch	CB
B label	unconditional branch	B
B.cond label	branch on flag (EQ,NE,LT,...)	CB
BL label	branch & link: $\text{LR} = \text{PC} + 4$, then branch	B
BR Xn	branch to address in Xn (return)	R

Conditionals

CMP Xn, Xm is SUBS XZR, Xn, Xm — it subtracts and sets the **N,Z,C,V** condition flags but discards the result. Then B.cond reads those flags.

7 C → Assembly, Calling Conventions & the Stack

7.1 Translating C control flow

if (a < b) { S }

```

SUBS XZR, X_a, X_b    // set flags
B.GE Lend             // skip if !(a<b)
...S...
Lend:

```

while (i != 0) { S; i--; }

```


Lp: CBZ X_i, Lend
...S...
SUBI X_i, X_i, #1
B Lp
Lend:

```

Arrays: `a[i]` for a 64-bit long array at base `Xa`: `LSL Xt,Xi,#3 (i×8)`, `ADD Xt,Xa,Xt`, `LDUR Xd,[Xt,#0]`. **Structs:** field at fixed byte offset from the struct's base pointer.

7.2 The stack and stack frames

The **stack** stores per-call data: saved registers, return address, locals, spilled args. It **grows toward lower addresses**. `SP` points at the current top.



The calling-convention contract

Caller-saved (X9–X15, X0–X7): if the *caller* needs a value after the call, the caller must save it first — the callee may clobber these freely.

Callee-saved (X19–X28, FP, LR): if the *callee* wants to use one, it must save the old value on entry and restore it before returning. The caller sees them unchanged.

Arguments: first 8 in X0–X7; extras on the stack. **Return value:** X0. **Return address:** BL puts it in LR; BR LR returns.

7.3 Standard function prologue / epilogue

Prologue (on entry):

```

SUBI SP, SP, #16    // make frame
STUR LR, [SP,#8]    // save return addr
STUR X19, [SP,#0]   // save callee reg
... body ...

```

Epilogue (before return):

```

LDUR X19, [SP,#0]   // restore
LDUR LR, [SP,#8]    // restore ret addr
ADDI SP, SP, #16    // pop frame
BR LR               // return

```

Common Mistake

For **recursive** or non-leaf functions you *must* save LR — otherwise the inner BL overwrites it and you can never return. Leaf functions that make no calls and use only temporaries may skip the frame entirely.

8 Digital Logic

8.1 Boolean algebra & gates

- $\text{AND} = A \cdot B$, $\text{OR} = A + B$, $\text{NOT} = \bar{A}$.
 - $\text{NAND} = \overline{A \cdot B}$, $\text{NOR} = \overline{A + B}$, $\text{XOR} = A \oplus B$ (1 iff inputs differ).
 - **NAND and NOR are universal** — any function can be built from only NANDs (or only NORs).
- Key identities**
- De Morgan: $\overline{A \cdot B} = \bar{A} + \bar{B}$, $\overline{A + B} = \bar{A} \cdot \bar{B}$.
 - Distributive: $A(B + C) = AB + AC$.
 - Absorption: $A + AB = A$.
 - $A + \bar{A} = 1$, $A \cdot \bar{A} = 0$.

8.2 Combinational building blocks

- **Multiplexer (MUX):** 2^n data inputs, n select lines, 1 output. Picks one input. A 2:1 MUX: $\text{out} = \bar{S} \cdot A + S \cdot B$. *The fundamental “choose” element in a datapath.*
- **Decoder:** n inputs $\rightarrow 2^n$ outputs, exactly one hot. Used to select registers / memory rows.
- **Adder:** full adder per bit, carry rippling. $\text{Sum} = A \oplus B \oplus C_{in}$, $C_{out} = AB + (A \oplus B)C_{in}$. Subtraction $= A + \bar{B} + 1$.

8.3 Sum-of-products from a truth table

1. Write the truth table for the function.
2. For every row where output = 1, write an AND term (a *minterm*) of the inputs (complemented where the input is 0).
3. OR all those minterms. Optionally simplify with Boolean algebra or a Karnaugh map.


8.4 Sequential logic: state & clocks

Combinational logic has no memory; **sequential** logic stores state.

- **D flip-flop:** on the clock edge, captures input D into output Q and holds it until the next edge. The atom of state (registers = arrays of D-FFs).
- **Clock period** must exceed the longest combinational path between flip-flops (plus setup time): $T_{clk} \geq t_{cq} + t_{logic} + t_{setup}$. This is what limits clock speed.

8.5 Finite State Machines (FSM)

A machine with state registers + combinational “next-state” and “output” logic.



To design one: (1) draw the state diagram, (2) assign binary codes to states, (3) build the truth table mapping (current state, input) $\rightarrow$ (next state, output), (4) derive logic via sum-of-products. Mealy outputs depend on state+input; Moore on state only.


9 Single-Cycle Datapath

Idea

Execute one entire instruction every clock cycle. Simple, but the clock must be slow enough for the *worst-case* instruction (usually `lw`, which touches PC, regs, ALU, *and* memory).

9.1 The five things every instruction does

IF fetch instruction → **ID** decode & read registers → **EX** ALU operation → **MEM** access data memory → **WB** write result to a register. Single-cycle does all five in one (long) cycle.



Control signals (set by the opcode) steer the MUXes: **RegWrite**, **ALUSrc** (register vs. sign-extended offset), **MemRead/MemWrite**, **MemToReg** (ALU result vs. loaded data), **Branch**, **PCSrc**.

10 Multi-Cycle Datapath

Idea

Break each instruction into several *short* cycles (roughly one per stage). A short clock period plus reusing hardware (one ALU, one memory) across cycles. Different instructions take different numbers of cycles (e.g. `lw`=5, `add`=4, `beq`=3).

Trade-off vs. single-cycle: clock period set by the *slowest stage*, not the slowest *instruction*; hardware reused; controlled by an FSM. But more cycles per instruction and more control complexity. This motivates **pipelining**: keep the short cycle, but overlap instructions so you finish (close to) one per cycle.


11 Pipelining

The laundry analogy

Don't wait for one load of laundry to wash→dry→fold before starting the next. Once the washer is free, start washing load 2 while load 1 dries. **Pipelining overlaps the stages of consecutive instructions.** Latency per instruction stays the same; *throughput* rises to ≈ 1 instruction/cycle.

11.1 The classic 5-stage pipeline

Stages **IF**, **ID**, **EX**, **MEM**, **WB** separated by **pipeline registers** (IF/ID, ID/EX, EX/MEM, MEM/WB) that carry each instruction's data and control forward.



In steady state (cycle 5+), **five instructions are in flight at once**, one per stage.

Pipeline speedup

Ideal speedup $\approx$ number of stages. For N instructions on a k -stage pipeline:


$$\text{cycles} = N + (k - 1) \quad (\text{the } k - 1 \text{ is the fill/drain "startup"}).$$

As $N \rightarrow \infty$, $\text{CPI} \rightarrow 1$. Real speedup is less due to **hazards**.

11.2 Hazards — the three reasons pipelines stall

1. **Structural hazard:** two instructions need the same hardware in the same cycle. Solved by duplicating resources (separate instruction & data memories; write regs in first half of a cycle, read in the second).
2. **Data hazard:** an instruction needs a result that an earlier, still-in-flight instruction hasn't written back yet.
3. **Control hazard:** after a branch, we don't yet know which instruction to fetch next.

Data hazards & forwarding



Forwarding (bypassing): route the ALU result directly from a later pipeline register (EX/MEM or MEM/WB) back into the EX stage's input, instead of waiting for WB. Eliminates most data-hazard stalls.

The load-use hazard cannot be fully forwarded

A **lw** produces its value at the end of **MEM**, but the next instruction needs it in **EX** — one cycle too early. You *must insert one stall (bubble)* (or the compiler reorders code). After the bubble, forward from MEM/WB.

Control hazards & branch handling

A **beq** resolves late (its outcome/target known in EX or MEM). Meanwhile the pipeline keeps fetching. Solutions:

- **Stall** until the branch resolves (simple, slow).
- **Predict not-taken:** keep fetching the next sequential instructions; if the branch *is* taken, **flush** (squash) the wrongly-fetched instructions (turn them into noops).
- **Branch prediction** (dynamic): a predictor (e.g. a 1- or 2-bit saturating counter per branch) guesses taken/not-taken from history; correct guesses cost nothing, mispredicts pay the flush penalty.
- **Resolve branches earlier** in the pipeline to shrink the penalty.

Stall/flush cost


Extra cycles = ($\#$ hazards) $\times$ (penalty per hazard). A branch resolved in stage s that was predicted wrong flushes $s - 1$ instructions $\Rightarrow$ penalty $s - 1$ cycles.

Dependency types (know the vocabulary)

- **RAW (Read-After-Write)** — a *true* dependence: an instruction reads a register a prior instruction writes. **This is the only hazard in the LC2K in-order pipeline**, and the one forwarding fixes.
- **WAR / WAW** — “false” (name) dependences that only matter with out-of-order or register renaming; **they do not occur** in the simple in-order 5-stage pipeline.

Dynamic branch prediction: the 2-bit saturating counter

A 1-bit predictor mispredicts *twice* per loop (on the last *and* first iteration). A **2-bit** predictor adds hysteresis — it must be wrong twice in a row to flip its prediction, so a loop branch mispredicts only once (on loop exit).



Predict **taken** in the two left states, **not-taken** in the two right.
T = branch actually taken, N = not taken.

11.3 The LC2K Pipeline Lab (Project 3) — implementation cheat sheet

What Project 3 asks for

Write a C simulator of the LC2K running on a **5-stage pipeline** with **full forwarding**, a **1-cycle load-use stall**, and **predict-not-taken** branches resolved in MEM. Every cycle you advance all five stages “simultaneously” and print the state. Correctness == your final register/memory values and total cycle count match the reference.

The pipeline-register and machine state structs

Five pipeline registers sit *between* the stages; each carries the `instr` plus whatever that boundary needs. Use the exact structs in the provided starter code — the fields below are the standard set:

```
#define NOOPINSTR 0x1c00000 // opcode 111 (noop) in bits 24-22

typedef struct IFIDStruct { int instr; int pcPlus1; } IFIDType;
typedef struct IDEXStruct { int instr; int pcPlus1;
                           int readRegA; int readRegB; int offset; } IDEXType;
typedef struct EXMEMStruct { int instr; int branchTarget;
                             int aluResult; int readRegB; } EXMEMType;
typedef struct MEMWBStruct { int instr; int writeData; } MEMWBType;
typedef struct WBENDStruct { int instr; int writeData; } WBENDType;

typedef struct stateStruct {
    int pc;
    int instrMem[NUMMEMORY]; int dataMem[NUMMEMORY];
    int reg[NUMREGS]; int numMemory;
    IFIDType IFID; IDEXType IDEX; EXMEMType EXMEM;
    MEMWBType MEMWB; WBENDType WBEND;
    int cycles; // cycles simulated so far
} stateType;
```

The double-buffering rule — this is THE structural trick

Keep two copies: `state` (this cycle’s values, read-only) and `newState` (next cycle, you write it). Each stage **reads from state and writes into newState**. At the end of the cycle do `state = newState`. **Iron rule:** state never

appears on the left-hand side of an assignment; `newState` never appears on the right. This makes all five stages update “at the same clock edge,” exactly like real hardware.

What each stage does (and which register it writes)

Stage	Work (reads state, writes <code>newState</code>)	Writes
IF	fetch <code>instrMem[pc]</code> ; <code>newState.pc = pc+1</code>	IFID
ID	decode <code>IFID.instr</code> ; read <code>reg[regA]</code> , <code>reg[regB]</code> ; sign-extend offset	IDEX
EX	compute: <code>add/nor</code> ALU op; <code>lw/sw</code> address <code>regA+offset</code> ; <code>branchTarget = IDEX.pcPlus1 + offset</code> ; apply forwarding here	EXMEM
MEM	<code>lw</code> : read <code>dataMem</code> ; <code>sw</code> : write <code>dataMem</code> ; resolve beq (taken? squash)	MEMWB
WB	write <code>reg[dest]</code> from <code>MEMWB.writeData</code>	WBEND

The extra **WBEND** register exists *only* so forwarding can reach an instruction three slots behind (the register-file write effectively lands a cycle “late” in this model).

Forwarding (done entirely in EX) — the exact algorithm

For the instruction now in EX, you need up-to-date `regA` and `regB`. Check the three older instructions ahead of it and take the value from the **most recent (closest) producer**:

Forwarding priority: EXMEM > MEMWB > WBEND

For each source register *r* that EX needs:


1. if EXMEM holds an `add/nor/lw` whose `dest = r` $\Rightarrow$ use `EXMEM.aluResult` (`lw`’s value isn’t ready yet — see stall below);
2. else if MEMWB `dest = r` $\Rightarrow$ use `MEMWB.writeData`;
3. else if WBEND `dest = r` $\Rightarrow$ use `WBEND.writeData`;
4. else use the value read from the register file (`IDEX.readRegA/B`).

Check newest first and stop — a later write must override an earlier one.

Computing “dest” depends on the opcode

`add/nor` write `destReg = bits 2–0`. `lw` writes `regB = bits 18–16`. `sw`, `beq`, `noop`, `halt` write **no register** — never forward from them. Compute each producer’s destination by reading *its* opcode, not the consumer’s.

The load-use stall (the one case forwarding can’t fix)




Detect in the EX stage: if the instruction in IDEX is a `lw` whose `dest (regB)` equals `regA` or `regB` of the instruction currently in IFID. **To stall one cycle:**

1. **Do not advance PC** (`newState.pc = state.pc`).
2. **Freeze IFID** (`newState.IFID = state.IFID`) — re-decode the stalled instruction.
3. **Inject a bubble** into IDEX: set `newState.IDEX.instr = NOOPINSTR`.

Next cycle the dependent instruction reaches EX while the `lw` is in MEMWB — forwarding now works.

Branch handling: predict-not-taken, resolve in MEM



1. **Predict not-taken:** IF just keeps fetching `pc+1` — no special action for branches in IF/ID/EX.
2. In **EX**, still compute `branchTarget = IDEX.pcPlus1 + offset` and carry it (plus whether `regA==regB`) into EXMEM.
3. In **MEM**, if the `beq` is *taken*, you mispredicted: set `newState.pc = EXMEM.branchTarget` and **squash the 3 younger instructions** by setting `newState.IFID.instr`, `newState.IDEX.instr`, and `newState.EXMEM.instr` to `NOOPINSTR`.

Penalty = 3 cycles per taken branch (since resolution is in MEM, stage 4, $\Rightarrow 4 - 1 = 3$). A not-taken branch costs nothing.

Per-cycle driver & start/stop

1. Init: `pc=0`, all five pipeline-register `instr` fields = `NOOPINSTR`, `cycles=0`.
2. Each iteration: `printState(&state)` *first*, then compute `WB→MEM→EX→ID→IF` into `newState` (any stage order is fine since you only read `state`), apply stall/squash, then `state = newState`, `state.cycles++`.
3. **Halt** when the instruction reaching the **MEM** (some specs say WB) stage is `halt`: print the final state and stop. Do *not* halt the moment you fetch it — let it flow down the pipe.

Project 3 bugs that cost the most points

1. Forwarding from a `sw/beq` (they write no register).
2. Wrong dest field (`add=bits 2-0` vs `lw=bits 18-16`).
3. Checking forwarding *oldest*-first so an older value wrongly wins — go newest (EXMEM) first.
4. Forgetting to also freeze `IFID` and hold `pc` on a load-use stall (you get a bubble but lose the instruction).
5. Off-by-one on `branchTarget` — it is `pcPlus1 + offset`, *not* `pc + offset`.
6. Halting on fetch instead of letting `halt` drain to the end.
7. Writing to `state` instead of `newState` (stages stop being simultaneous).

12 Performance Analysis

The Iron Law of performance

$$\text{CPU time} = \underbrace{\text{Instructions}}_{\text{program/ISA/compiler}} \times \underbrace{\text{CPI}}_{\text{architecture}} \times \underbrace{\text{Clock cycle time}}_{\text{technology}} = \frac{\text{Instructions} \times \text{CPI}}{\text{Clock rate}}$$

- **CPI** = cycles per instruction (average). **Average CPI** = $\sum_i (\text{freq}_i \times \text{CPI}_i)$.
- **Speedup** = $\frac{\text{time}_{\text{old}}}{\text{time}_{\text{new}}}$. (> 1 means faster.)
- **MIPS** = $\frac{\text{Instruction count}}{\text{exec time} \times 10^6}$ (a flawed metric across ISAs).

Amdahl's Law — the “speed limit” of optimization

If a fraction f of execution is sped up by factor s :

$$\text{Speedup}_{\text{overall}} = \frac{1}{(1 - f) + \frac{f}{s}}$$

As $s \rightarrow \infty$, speedup $\rightarrow \frac{1}{1-f}$. **Make the common case fast:** optimizing a tiny fraction f can never help more than $1/(1-f)$.

Worked example: 40% of a program is parallelizable ($f=0.4$) and you speed that part up $4\times$ ($s=4$). Overall speedup $= \frac{1}{0.6+0.4/4} = \frac{1}{0.7} \approx 1.43\times$. Even with $s = \infty$ the cap is $1/0.6 \approx 1.67\times$.

13 Caches & the Memory Hierarchy

Why caches exist

DRAM is huge but slow (~ 100 s of cycles); registers are fast but tiny. A **cache** is a small fast SRAM holding recently/likely-used data. It works because of **locality**:

Temporal: data used now is likely used again soon (loops, variables).

Spatial: data near recently-used data is likely used soon (arrays, code). $\Rightarrow$ fetch a whole **block** (line) at a time.

13.1 Address breakdown

Every address is split into three fields (sizes depend on cache geometry):

identifies the block	picks the set/line	picks byte/word in block
TAG	INDEX	BLOCK OFFSET

The bit-counting recipe (do this every cache problem)

Let block size $= 2^b$ bytes, number of sets $= 2^s$, address $= m$ bits.

$$\text{offset bits} = b = \log_2(\text{block size}), \quad \text{index bits} = s = \log_2(\# \text{ sets}),$$


$$\text{tag bits} = m - s - b, \quad \# \text{sets} = \frac{\text{cache size}}{(\text{block size}) \times (\text{associativity})}.$$

13.2 Mapping: direct, set-associative, fully associative


Type	A block can go...	# sets / ways
Direct-mapped	in exactly 1 line (set)	sets = lines, 1 way
n -way set-associative	in any of n lines of its set	sets = lines/ n
Fully associative	in any line	1 set, all ways

More associativity $\Rightarrow$ fewer conflict misses but more comparators and slower hits.

Direct-mapped



2-way set-associative



13.3 The 3 C's — classify every miss

- **Compulsory (cold):** first-ever access to a block. Unavoidable except by prefetching / bigger blocks.

- **Capacity:** cache simply can't hold the working set. Fixed only by a bigger cache.
- **Conflict:** block was evicted because too many blocks map to the same set (would have hit in a fully-associative cache). Fixed by more associativity.

13.4 Replacement & write policies

Replacement (which line to evict in an associative set): **LRU** (least recently used), **FIFO**, or **random**.

On a write hit: *Write-through* (update cache *and* memory; simple, more traffic — usually with a write buffer) vs. *Write-back* (update cache only, set a **dirty bit**; write to memory on eviction; less traffic).

On a write miss: *Write-allocate* (load block into cache, then write) — pairs with write-back; vs. *No-write-allocate* (write straight to memory) — pairs with write-through.

13.5 Performance: AMAT

Average Memory Access Time

$$\text{AMAT} = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$$

For multi-level caches, recurse on the miss penalty:

$$\text{AMAT} = HT_{L1} + MR_{L1}(HT_{L2} + MR_{L2} \times MP_{L2})$$

Local miss rate = misses in this level / accesses to this level. **Global miss rate** = misses in this level / *total* CPU accesses.

Worked example: hit time 1 cycle, miss rate 5%, miss penalty 100 cycles. $\text{AMAT} = 1 + 0.05 \times 100 = 6$ cycles. Cut miss rate to 2%: $1 + 0.02 \times 100 = 3$ cycles ($2\times$ faster memory access).

Cache-aware programming

Loop over arrays in **memory order** (row-major in C: vary the *last* index in the inner loop) to exploit spatial locality. Blocking/tiling keeps a working set inside the cache to exploit temporal locality. This is the “performance optimization” the course rewards.


14 Virtual Memory

Why virtual memory

Give every process its own private, contiguous-looking address space that's larger than physical RAM, with **protection** (processes can't touch each other's memory) and **relocation** (programs don't care where they physically land). The OS + hardware translate **virtual addresses** → **physical addresses**, paging data to/from disk on demand.

14.1 Paging and address translation

Memory is divided into fixed-size **pages** (virtual) and **frames** (physical), e.g. 4KB. A virtual address splits into a **Virtual Page Number (VPN)** + **page offset**; the offset is unchanged by translation.



Translation arithmetic

$$\text{offset bits} = \log_2(\text{page size}), \quad \# \text{virtual pages} = \frac{\text{virtual addr space}}{\text{page size}},$$

$$\text{PFN bits} = \log_2\left(\frac{\text{physical mem}}{\text{page size}}\right), \quad \text{Physical addr} = \text{PFN} \parallel \text{offset}.$$


14.2 Page table

A per-process table mapping $\text{VPN} \rightarrow \text{PFN}$, with **valid**, **dirty**, **protection** bits. Lives in memory. If **valid**=0 on access $\Rightarrow$ **page fault**: OS fetches the page from disk (very slow), updates the table, resumes.

Multi-level page tables: a single flat table for a 64-bit space would be enormous, so use a tree — the VPN is split into several indices, each indexing the next-level table. Only tables for in-use regions exist, saving space (at the cost of multiple memory lookups per translation).

14.3 TLB — making translation fast**Key Idea**

A naive load now costs *two* memory accesses (page table, then data). The **TLB (Translation Lookaside Buffer)** is a small fast cache of recent $\text{VPN} \rightarrow \text{PFN}$ translations. TLB hit $\Rightarrow$ translate in ~ 1 cycle.



Full access sequence for a load: (1) check **TLB** for VPN. (2) TLB miss $\rightarrow$ walk the **page table** (and on **valid**=0, take a **page fault** to disk). (3) form the physical address. (4) look it up in the **cache**; on a cache miss, go to memory. So three nested “maybe-miss” structures: **TLB** $\rightarrow$ **page table** $\rightarrow$ **cache**.

15 Rapid-Fire Exam Checklists**LC2K assembly problem**

1. Remember **r0**=0 by convention. 2. No **addi**/immediate-load — use **.fill** + **lw**. 3. **beq** offset = targetAddr – (beqAddr+1). 4. Negate via **nor**+add 1. 5. Each line = one word = one address.

Pipeline / hazard problem

1. Draw the cycle-by-cycle diagram. 2. Mark every RAW dependency between in-flight instructions. 3. Can forwarding fix it? (Yes, except load-use $\rightarrow$ 1 stall.) 4. Branch: add flush penalty = (resolve stage – 1). 5. Total cycles = $N + (k - 1) + \text{stalls} + \text{flushes}$.

Cache problem

1. offset = $\log_2(\text{block})$, index = $\log_2(\# \text{ sets})$, tag = rest. 2. $\# \text{ sets} = \text{size} / (\text{block} \times \text{assoc})$. 3. Trace each access; classify miss as compulsory/capacity/conflict. 4. AMAT = hit + missrate $\times$ penalty. 5. Don’t forget **valid/dirty/tag overhead bits**.

Virtual memory problem

1. $\text{offset} = \log_2(\text{page size})$. 2. $\text{VPN bits} = \text{virtAddrBits} - \text{offset}$. 3. PFN bits from physical size. 4. Page-table entries = $2^{\text{VPN bits}}$ per level. 5. Access path: TLB $\rightarrow$ page table $\rightarrow$ cache $\rightarrow$ memory.

End of cheat sheet — Go Blue. Always re-verify exact opcodes / bit ranges and the current term's policies against the official EECS 370 spec before your exam.